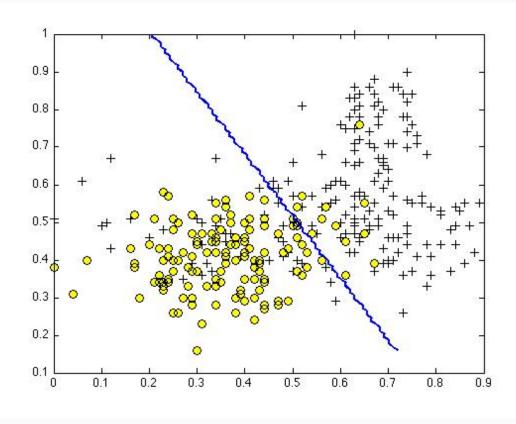
You submitted this quiz on **Mon 5 May 2014 8:41 AM IST**. You got a score of **5.00** out of **5.00**.

Question 1

Suppose you have trained an SVM classifier with a Gaussian kernel, and it learned the following decision boundary on the training set:



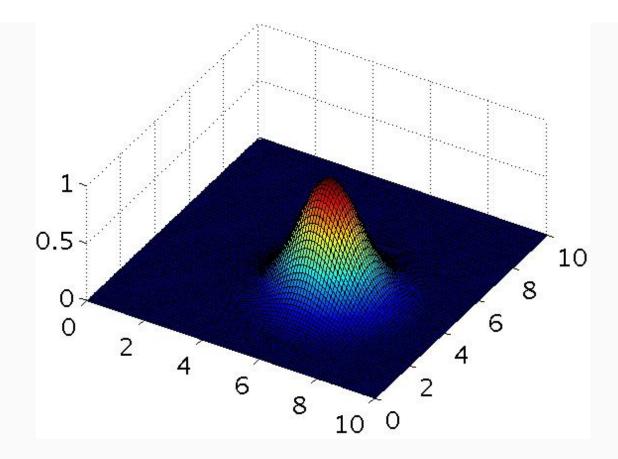
You suspect that the SVM is underfitting your dataset. Should you try increasing or decreasing C? Increasing or decreasing σ^2 ?

Your Answer	Score	Explanation
It would be reasonable to try increasing C . It would also be		
reasonable to try increasing σ^2 .		
It would be		

reasonable to try decreasing C . It would also be reasonable to try increasing σ^2 .		
It would be reasonable to try increasing C . It would also be reasonable to try decreasing σ^2 .	✓ 1.00	The figure shows a decision boundary that is underfit to the training set, so we'd like to lower the bias / increase the variance of the SVM. We can do so by either increasing the parameter C or decreasing σ^2 .
It would be reasonable to try decreasing C . It would also be reasonable to try decreasing σ^2 .		
Total	1.00 / 1.00	

Question 2

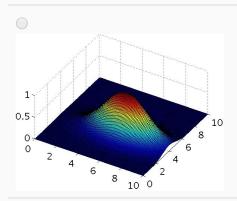
The formula for the Gaussian kernel is given by $\mathrm{similarity}(x,l^{(1)}) = \exp{(-\frac{||x-l^{(1)}||^2}{2\sigma^2})}$. The figure below shows a plot of $f_1 = \mathrm{similarity}(x,l^{(1)})$ when $\sigma^2 = 1$.

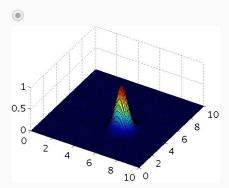


Which of the following is a plot of f_1 when $\sigma^2=0.25$?

Your Answer

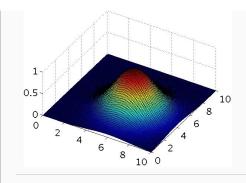
Explanation Score

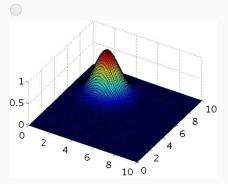




1.00

This figure shows a "narrower" Gaussian kernel centered at the same location which is the effect of decreasing σ^2 .

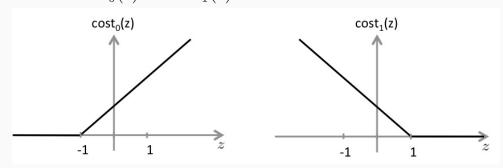




Total 1.00 / 1.00

Question 3

The SVM solves $\min_{\theta} \ C \sum_{i=1}^m y^{(i)} \mathrm{cost}_1(\theta^T x^{(i)}) + (1-y^{(i)}) \mathrm{cost}_0(\theta^T x^{(i)}) + \sum_{j=1}^n \theta_j^2$ where the functions $\mathrm{cost}_0(z)$ and $\mathrm{cost}_1(z)$ look like this:



The first term in the objective is: $C\sum_{i=1}^m y^{(i)} \mathrm{cost}_1(\theta^T x^{(i)}) + (1-y^{(i)}) \mathrm{cost}_0(\theta^T x^{(i)})$. This first term will be zero if two of the following four conditions hold true. Which are the two conditions that would guarantee that this term equals zero?

Your Answer		Score	Explanation
For every example with $y^{(i)}=1$, we have that $ heta^T x^{(i)} \geq 1$.	~	0.25	For examples with $y^{(i)}=1$, only the $\mathrm{cost}_1(\theta^Tx^{(i)})$ term is present. As you can see in the graph, this will be zero for all inputs greater than or equal to 1.
	~	0.25	For examples with $y^{(i)} = 0$, only the $\mathrm{cost}_0(heta^T x^{(i)})$

with $y^{(i)}=0$, we have that $ heta^T x^{(i)} \leq -1$.		term is present. As you can see in the graph, this will be zero for all inputs less than or equal to -1.
\Box For every example with $y^{(i)}=0$, we have that $ heta^T x^{(i)} \leq 0$.	✔ 0.25	$\cot_0(\theta^T x^{(i)})$ is still non-zero for inputs between -1 and 0, so being less than or equal to 0 is insufficient.
\Box For every example with $y^{(i)}=1$, we have that $ heta^T x^{(i)} \geq 0$.	✔ 0.25	$\cot_1(\theta^T x^{(i)})$ is still non-zero for inputs between 0 and 1, so being greater than or equal to 0 is insufficient.
Total	1.00 / 1.00	

Question 4

Suppose you have a dataset with n = 10 features and m = 5000 examples. After training your logistic regression classifier with gradient descent, you find that it has underfit the training set and does not achieve the desired performance on the training or cross validation sets. Which of the following might be promising steps to take? Check all that apply.

Your Answer	Score	Explanation
Use a different optimization method since using gradient descent to train logistic regression might result in a local minimum.	✓ 0.25	The logistic regression cost function is convex, so gradient descent will always find the global minimum.
	✔ 0.25	When you add more features, you increase the variance of your model, reducing the chances of underfitting.
	✔ 0.25	A neural network with many hidden units is a more complex (higher variance) model than logistic regression, so it is less likely to underfit the data.
□ Increase the regularization parameter λ .	✔ 0.25	You are already underfitting the data, and increasing the regularization parameter only makes underfitting stronger.
Total	1.00 / 1.00	

Question 5

Which of the following statements are true? Check all that apply.

Your Answer	Score	Explanation
☐ If you are training multi- class SVMs with the one- vs-all method, it is not possible to use a kernel.	✔ 0.25	Each SVM you train in the one-vs-all method is a standard SVM, so you are free to use a kernel.
Suppose you are using SVMs to do multi-class classification and would like to use the one-vs-all approach. If you have K different classes, you will train K - 1 different SVMs.	✔ 0.25	The one-vs-all method requires that we have a separate classifier for every class, so you will train K different SVMs.
The maximum value of the Gaussian kernel (i.e., $sim(x, l^{(1)})$) is 1.	✔ 0.25	When $x=l^{(1)}$, the Gaussian kernel has value $\exp\left(0\right)=1$, and it is less than 1 otherwise.
${f V}$ Suppose you have 2D input examples (ie, $x^{(i)} \in {\Bbb R}^2$). The decision boundary of the SVM (with the linear kernel) is a straight line.	✔ 0.25	The SVM without any kernel (ie, the linear kernel) predicts output based only on $\theta^T x$, so it gives a linear / straight-line decision boundary, just as logistic regression does.
Total	1.00 / 1.00	