You submitted this quiz on **Sun 13 Apr 2014 2:16 PM IST**. You got a score of **4.00** out of **5.00**. You can attempt again in 10 minutes.

Question 1

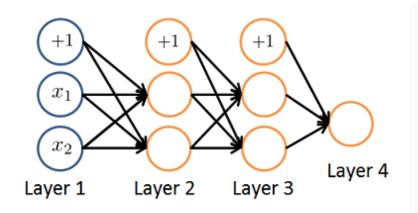
Consider the following neural network which takes two binary-valued inputs $x_1, x_2 \in \{0, 1\}$ and outputs $h_{\Theta}(x)$. Which of the following logical functions does it (approximately) compute?



Your Answer	Score	Explanation
○ NAND (meaning "NOT AND")		
○ OR		
AND		
XOR (exclusive OR)	x 0.00	You cannot approximate the XOR function with a two layer neural network.
Total	0.00 / 1.00	

Question 2

Consider the neural network given below. Which of the following equations correctly computes the activation $a_1^{(3)}$? Note: g(z) is the sigmoid activation function.



four Answer	Score	Explanation

$$a_1^{(3)} = g(\Theta_{1,0}^{(1)}a_0^{(2)} + \Theta_{1,1}^{(1)}a_1^{(2)} + \Theta_{1,2}^{(1)}a_2^{(2)})$$

$$a_1^{(3)} = g(\Theta_{1,0}^{(1)}a_0^{(1)} + \Theta_{1,1}^{(1)}a_1^{(1)} + \Theta_{1,2}^{(1)}a_2^{(1)})$$

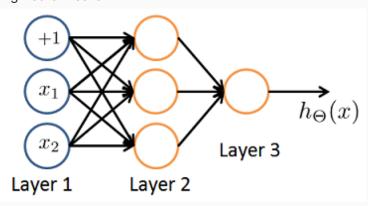
$$a_1^{(3)} = g(\Theta_{1,0}^{(2)}a_0^{(2)} + \Theta_{1,1}^{(2)}a_1^{(2)} + \Theta_{1,2}^{(2)}a_2^{(2)}) \qquad \qquad \text{This correctly uses the first row of } \Theta^{(2)} \text{ and includes the "+1" term of } a_0^{(2)}.$$

$$a_1^{(3)} = g(\Theta_{2,0}^{(2)}a_0^{(2)} + \Theta_{2,1}^{(2)}a_1^{(2)} + \Theta_{2,2}^{(2)}a_2^{(2)})$$

Total 1.00 / 1.00

Question 3

You have the following neural network:



You'd like to compute the activations of the hidden layer $a^{(2)} \in \mathbb{R}^3$. One way to do so is the

following Octave code:

```
% Theta1 is Theta with superscript "(1)" from lecture
% ie, the matrix of parameters for the mapping from layer 1 (input) to layer 2
% Theta1 has size 3x3
% Assume 'sigmoid' is a built-in function to compute 1 / (1 + exp(-z))

a2 = zeros (3, 1);
for i = 1:3
    for j = 1:3
        a2(i) = a2(i) + x(j) * Theta1(i, j);
    end
        a2(i) = sigmoid (a2(i));
end
```

You want to have a vectorized implementation of this (i.e., one that does not use for loops). Which of the following implementations correctly compute $a^{(2)}$? Check all that apply.

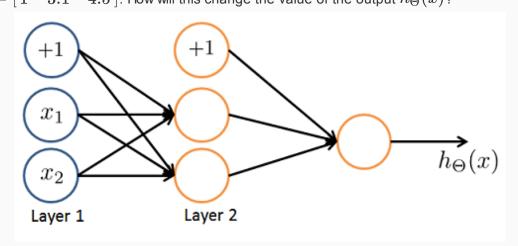
Your Answer		Score	Explanation
z = Theta1 * x; a2 = sigmoid (z);	~	0.25	This version computes $a^{(2)}=g(\Theta^{(1)}x)$ correctly in two steps, first the multiplication and then the sigmoid activation.
a2 = sigmoid (x * Theta1);	~	0.25	The order of the multiplication is important, this will not work as x is a vector of size 3×1 while Theta1 is a matrix of size $3x3$.
a2 = sigmoid (T heta2 * x);	~	0.25	$\Theta^{(2)}$ specifies the parameters from the second to third layers, not first to second.
z = sigmoid(x); a2 = Theta1 * z ;	~	0.25	You should apply the sigmoid function after multiplying with $\Theta^{(1)}$, not before.
Total		1.00 / 1.00	

Question 4

You are using the neural network pictured below and have learned the parameters

$$\Theta^{(1)}=egin{bmatrix}1&2.1&1.3\\1&0.6&-1.2\end{bmatrix}$$
 (used to compute $a^{(2)}$) and $\Theta^{(2)}=[1&4.5&3.1\,]$ (used to

compute $a^{(3)}$ } as a function of $a^{(2)}$). Suppose you swap the parameters for the first hidden layer between its two units so $\Theta^{(1)}=\begin{bmatrix}1&0.6&-1.2\\1&2.1&1.3\end{bmatrix}$ and also swap the output layer so $\Theta^{(2)}=\begin{bmatrix}1&3.1&4.5\end{bmatrix}$. How will this change the value of the output $h_{\Theta}(x)$?



Your Answer	Score	Explanation
It will increase.		
It will decrease		
 Insufficient information to tell: it may increase or decrease. 		
It will stay the same.	✓ 1.00	Swapping $\Theta^{(1)}$ swaps the hidden layers output a $^{(2)}$. But the swap of $\Theta^{(2)}$ cancels out the change, so the output will remain unchanged.
Total	1.00 / 1.00	

Question 5

Which of the following statements are true? Check all that apply.

Your Answer		Score	Explanation
A two layer (one input layer, one output layer; no hidden layer) neural network can represent the XOR function.	~	0.25	We must compose multiple logical operations by using a hidden layer to represent the XOR function.

If a neural network is overfitting the data, one solution would be to decrease the regularization parameter λ .	~	0.25	A smaller value of λ allows the model to more closely fit the training data, thereby increasing the chances of overfitting.
Any logical function over binary-valued (0 or 1) inputs x_1 and x_2 can be (approximately) represented using some neural network.	•	0.25	Since we can build the basic AND, OR, and NOT functions with a two layer network, we can (approximately) represent any logical function by composing these basic functions over multiple layers.
If a neural network is overfitting the data, one solution would be to increase the regularization parameter λ .	~	0.25	A larger value of λ will shrink the magnitude of the parameters Θ , thereby reducing the chance of overfitting the data.
Total		1.00 /	