Feedback — VI. Logistic Regression

Help

You submitted this quiz on **Tue 1 Apr 2014 7:43 PM IST**. You got a score of **5.00** out of **5.00**.

Question 1

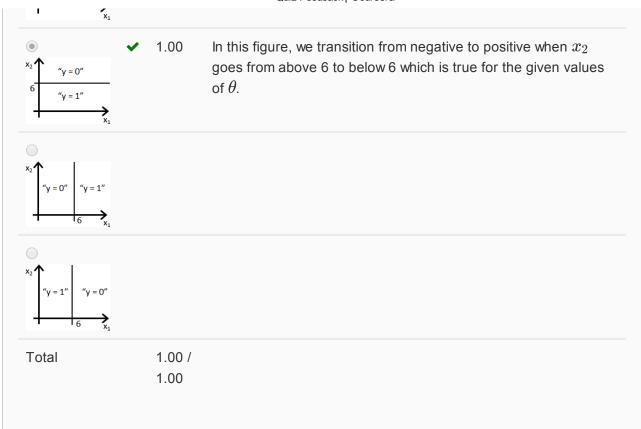
Suppose that you have trained a logistic regression classifier, and it outputs on a new example x a prediction $h_{\theta}(x)$ = 0.4. This means (check all that apply):

Your Answer		Score	Explanation
lacksquare Our estimate for $P(y=1 x; heta)$ is 0.6.	~	0.25	$h_{ heta}(x)$ gives $P(y=1 x; heta)$, not $1-P(y=1 x; heta)$.
\square Our estimate for $P(y=0 x; heta)$ is 0.4.	~	0.25	$h_{ heta}(x)$ is $P(y=1 x; heta)$, not $P(y=0 x; heta)$
ightharpoonup Our estimate for $P(y=1 x; heta)$ is 0.4.	~	0.25	$h_{ heta}(x)$ is precisely $P(y=1 x; heta)$, so each is 0.4.
ightharpoonup Our estimate for $P(y=0 x; heta)$ is 0.6.	~	0.25	Since we must have $P(y=0 x;\theta)=1-P(y=1 x;\theta) \text{, the former is 1 - 0.4 = 0.6}.$
Total		1.00 / 1.00	

Question 2

Suppose you train a logistic classifier $h_{\theta}(x)=g(\theta_0+\theta_1x_1+\theta_2x_2)$. Suppose $\theta_0=6, \theta_1=0, \theta_2=-1$. Which of the following figures represents the decision boundary found by your classifier?

Your	Score	Explanation	
Answer			
0			
x₂			
6 " 0"			

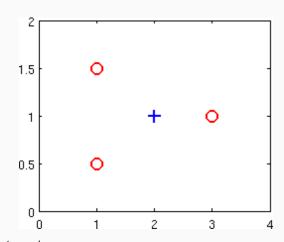


Question 3

Suppose you have the following training set, and fit a logistic regression classifier

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2).$$

x_1	x_2	y
1	0.5	0
1	1.5	0
2	1	1
3	1	0



Which of the following are true? Check all that apply.

Your Answer		Score	Explanation	
$ \checkmark \!$	~	0.25	The cost function $J(\theta)$ is always no negative for logistic regression.	
$lacksquare$ If we train gradient descent for enough iterations, for some examples $x^{(i)}$ in the training set it is possible to obtain $h_{ heta}(x^{(i)})>1.$	•	0.25	The function $g(z)$ the hypothesis $h_{\theta}(x)$ is the	in

Quiz Feedback Coursera		$\frac{1}{1+e^{-z}}$ which always lies between 0 and 1.
Adding polynomial features (e.g., instead using $h_{\theta}(x)=g(\theta_0+\theta_1x_1+\theta_2x_2+\theta_3x_1^2+\theta_4x_1x_2+\theta_5x_2^2)$) could increase how well we can fit the training data.	✔ 0.25	Adding new feature can only improve the fit on the training set: since setting $\theta_3 = \theta_4 = \theta_5 = 0$ makes the hypothesis the same as the original one, gradient descent with use those features (by making the corresponding θ_j non-zero) only if doing so improves the training set fit.
■ Because the positive and negative examples cannot be separated using a straight line, linear regression will perform as well as logistic regression on this data.	✓ 0.25	While it is true they cannot be separated, logistic regression will outperform linear regression since its cost function focuses on classification, not prediction.
Total	1.00 / 1.00	

Question 4

For logistic regression, the gradient is given by $rac{\partial}{\partial heta_j} J(heta) = \sum_{i=1}^m (h_{ heta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$.

Which of these is a correct gradient descent update for logistic regression with a learning rate of α ? Check all that apply.

Your Answer		Score	Explanation
$oldsymbol{ heta} heta_j := heta_j - lpha rac{1}{m} \sum_{i=1}^m igg(rac{1}{1 + e^{- heta^T x^{(i)}}} - y^{(i)}igg) x_j^{(i)}$	~	0.25	This substitutes the exact form of $h_{ heta}(x^{(i)})$ used by logistic regression into the

(simultaneously update for all j).		gradient descent update.
$egin{aligned} & egin{aligned} & eta_j := heta_j - lpha rac{1}{m} \sum_{i=1}^m ig(h_ heta(x^{(i)}) - y^{(i)}ig) x^{(i)} \ & ext{(simultaneously update for all } j). \end{aligned}$	✓ 0.25	This incorrectly multiplies by the vector $\boldsymbol{x}^{(i)}$ in the summation rather than just $\boldsymbol{x}_j^{(i)}$.
$lacksquare heta := heta - lpha rac{1}{m} \sum_{i=1}^m \Big(heta^T x - y^{(i)} \Big) x^{(i)}.$	✓ 0.25	This vectorized version uses the linear regression hypothesis $\theta^T x$ instead of that for logistic regression.
extstyle hinspace hinspacee hinspace hinspace hinspace hinspace hinspace hinspace hi	✔ 0.25	This is a direct substitution of $\frac{\partial}{\partial \theta_j} J(\theta)$ into the gradient descent update.
Total	1.00	I

Question 5

Which of the following statements are true? Check all that apply.

Your Answer		Score	Explanation
$ \begin{tabular}{ll} \hline \mathscr{N} The cost \\ function $J(\theta)$ for \\ logistic regression \\ trained with $m \geq 1$ \\ examples is always \\ greater than or \\ equal to zero. \\ \end{tabular} $	~	0.25	The cost for any example $x^{(i)}$ is always ≥ 0 since it is the negative log of a quantity less than one. The cost function $J(\theta)$ is a summation over the cost for each eample, so the cost function itself must be greater than or equal to zero.
Linear regression always works well for classification if you classify by using a threshold on the prediction made by linear regression.	~	0.25	As demonstrated in the lecture, linear regression often classifies poorly since its training prodcedure focuses on predicting real-valued outputs, not classification.
For logistic regression, sometimes gradient descent will	~	0.25	The cost function for logistic regression is convex, so gradient descent will always converge to the global minimum. We still might use a more advanded optimization algorithm since they can be faster and don't require you to

		Quiz Feedback Coursera
converge to a local minimum (and fail to find the global minimum). This is the reason we prefer more advanced optimization algorithms such as fminunc (conjugate gradient/BFGS/L-BFGS/etc).		select a learning rate.
The one-vs-all technique allows you to use logistic regression for problems in which each $y^{(i)}$ comes from a fixed, discrete set of values.	✔ 0.25	If each $y^{(i)}$ is one of k different values, we can give a label to each $y^{(i)} \in \{1,2,\dots,k\}$ and use one-vs-all as described in the lecture.
Total	1.00 / 1.00	